Customer Acquisition Funnel Analysis

Prepared by: Sebastian Sava

Marketing performance case study

THE SITUATION

More leads and higher spend did not translate into more bookings

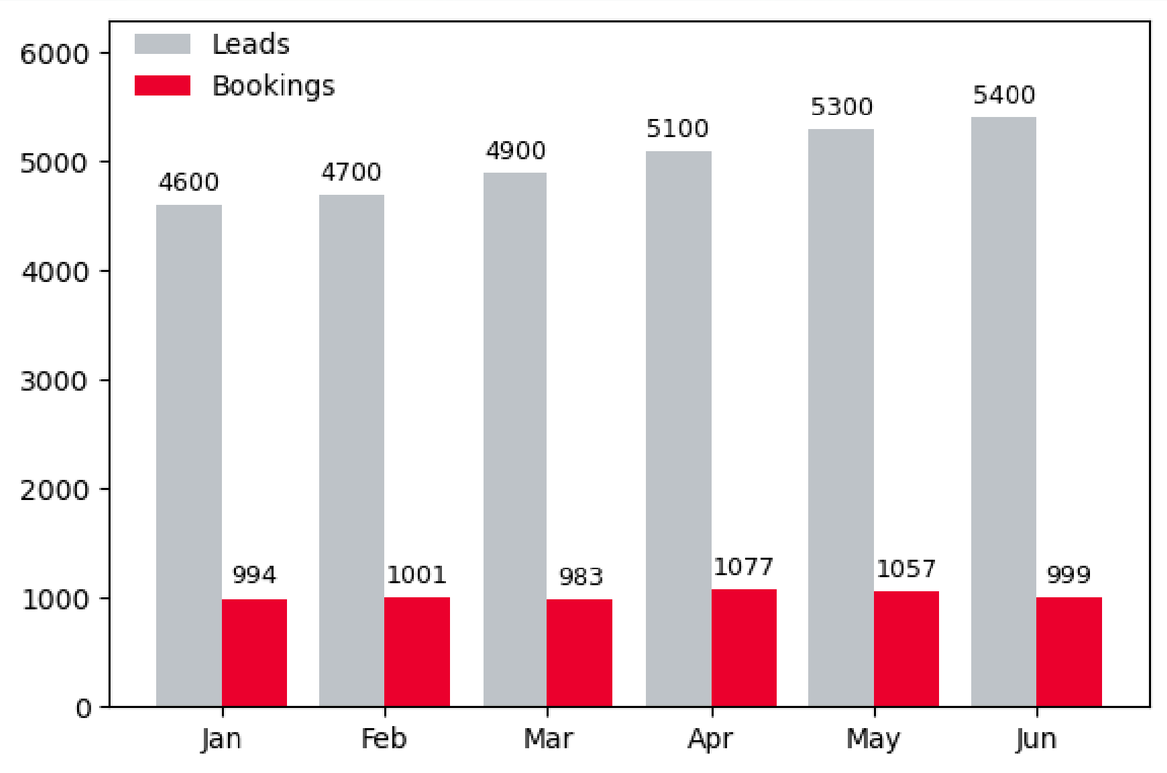
Across six months, 30,000 leads entered the process. 28,591 were called at least once, 19,431 spoke to an agent, and 6,111 booked a visit. Overall lead-to-booking conversion was 20.4%, with an average cost of EUR 291 per booking.

Lead growth did not translate into booking growth

Monthly lead volume rose from 4,600 in January to 5,400 in June, and spend rose with it, from 265k to 326k EUR. Bookings were 994 in January and 999 in June.

The cost of a booking rose 22%

The cost per lead changed little, from 58 to 60 EUR, so the increase in cost per booking is tied to the drop in conversion rather than higher lead prices. Conversion was 21.6% in January and 18.5% in June. The next slide examines whether that difference holds up.



The decline persists after several checks, with most of it occurring after connection

The decline remains after validation checks

Restricting the sample to leads with a full 13-day conversion window leaves June conversion at **18.6%**, compared with **18.5%** overall. Reweighting the channel mix to match January also does not materially change the result. The 3.1 percentage-point decline therefore remains after these checks.

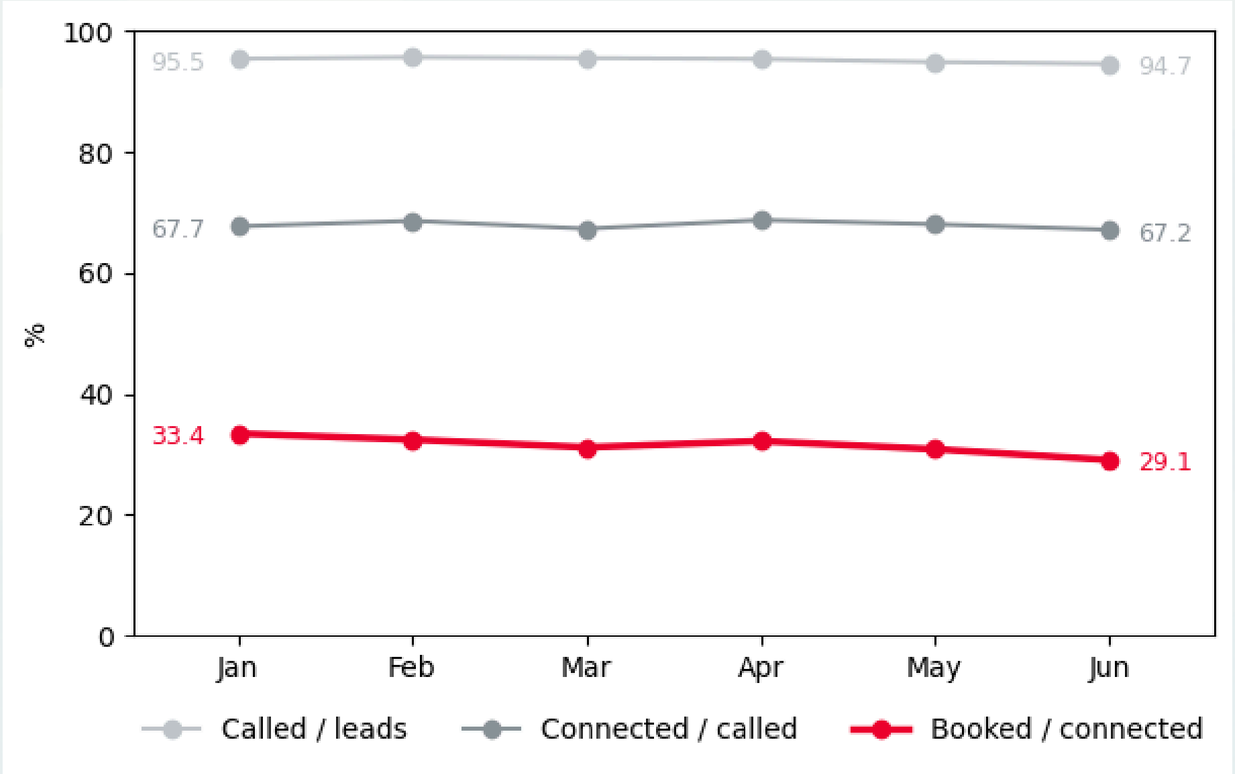
The loss is in the last stage of the funnel

Call coverage stayed around **95%** and the connection rate around **67%**. After connection, bookings per connected lead fell from 33.4% to **29.1%**, accounting for roughly 90% of the overall conversion decline. Channel A shows the same pattern, falling from 34.3% to 27.0%.

Cost per lead can be misleading

Channel C has the lowest cost per lead at EUR **44**, but its lower conversion rate results in a cost per booking of **EUR 304,** compared with **EUR 291** overall. Channel A accounts for 36% of spend and has a cost per booking of EUR 334.

Funnel stages over the six months (%)



Later attempts still produce bookings, but many leads stop early

Attempt count is affected by selection

Leads with only one recorded call converted at **27.9%**, compared with **11.8%** among leads that reached five calls. But this comparison is selected by outcome: once a lead books, calls stop.
Looking instead at each call, the observed booking rate is **11.9%** on the first attempt and **11.8%** on the fifth. These rates describe what happened; they do not estimate the causal value of adding another call.

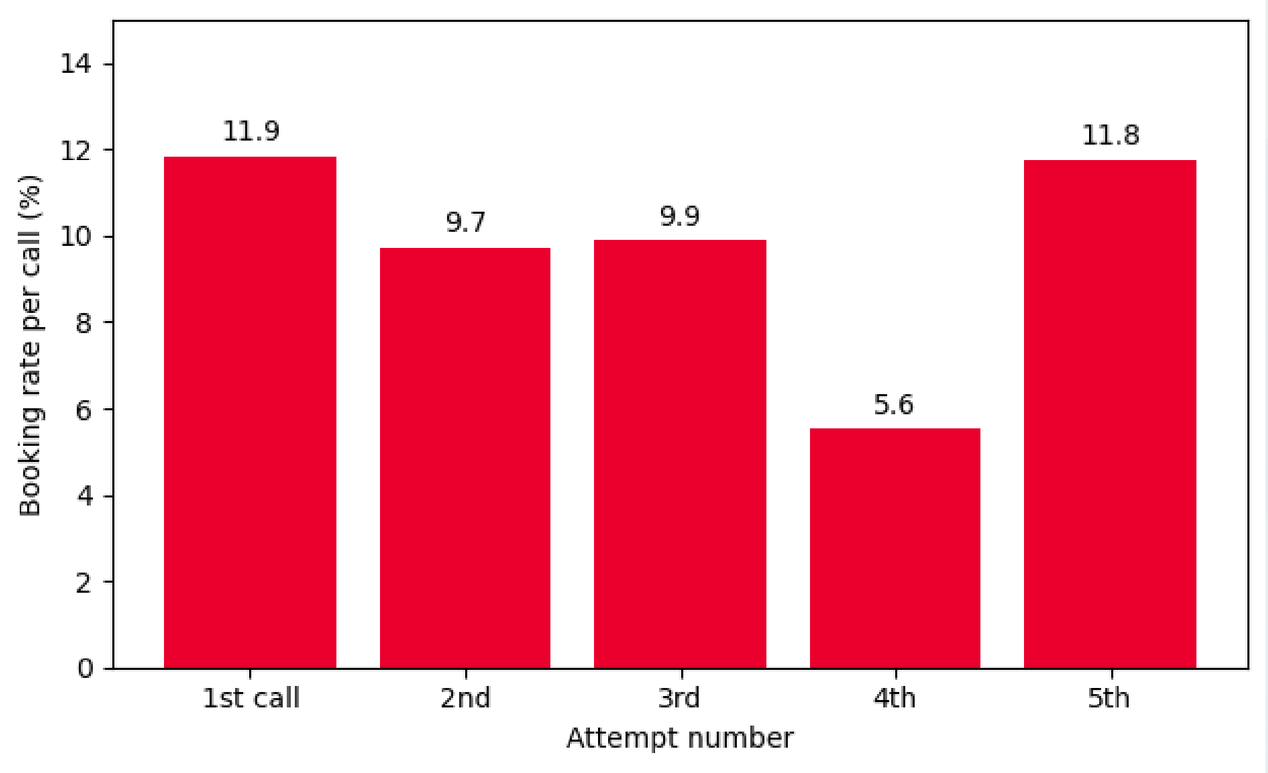
Most leads stop after one or two calls

21,617 unbooked leads stopped before a fifth attempt, including **8,730** after one call. After excluding leads recorded as not interested or unreachable, around **11,000** remain potentially relevant for further follow-up. Of these, **2,386** were last recorded as “Follow-up required.”

Previous outcome and next-call result

When another call did occur, the next call booked **16.0%** of the time after “Follow-up required”, compared with **8.4%** after “No answer”. This suggests previous call outcome may be useful for prioritisation, but these are observed rates, not causal estimates.

Observed booking rate per call, by attempt (%)



RECOMMENDATIONS | TASK 3

What to do next: improve follow-up, test routing, and measure beyond booking

LEAD TREATMENT

Follow up the leads who asked to be called back

2,386 unbooked leads were last marked as “Follow-up required” and received no further call. When similar leads did get another call, 16.0% booked versus 8.4% after “No answer.”

I would test one more attempt with these leads and measure how many extra bookings this generates per 100 additional calls.

OUTBOUND OPERATIONS

Decide who to call first when capacity is tight

Every channel gets about 1.9 attempts per lead, though channels differ a lot in how often their leads book. The last call outcome varies more than the channel does.

Effort is spread evenly while the return is not. Ordering the queue on outcome first and channel second is worth testing against current practice.

OUTBOUND COVERAGE

Find out why some leads are never called

1,409 leads were never called, increasing from 207 in January to 288 in June. Another 8,610 unbooked leads ended with a “No answer” and received fewer than five attempts.

The data does not tell us why these leads were not followed up further. If the team has available capacity, I would start with the leads that were never called.

INVESTMENT & MEASUREMENT

Test media changes rather than moving budget

Channel A costs more per booking than average, while B, E and F look better. But the costs here are synthetic and the data stops at the booked visit.

I would run a small test in one channel, and extend reporting to whether the visit happened, whether it led to a customer, and how long they stay.

The first two recommendations need no additional media spend, but they consume agent time that this dataset does not measure.

Assumptions, limitations and other considerations

What I checked in the data

- Coverage, keys, missing values and timestamp order all came back clean: complete cost grid, unique IDs, no orphan or missing values, no event out of order.
- The outbound table had 180 rows that were exact copies of another row. I removed them first. The 23 bookings that looked duplicated were inside those rows, so once gone the booking flags came to 6,111, matching the prospects table.
- Campaign_Name is the attempt number under a different name. Talk time never matches call length. Connected_Flag contains the same information as Contact_Outcome in this dataset.
- The dataset documentation states that costs and contact timing were randomised, so channel cost comparisons should be validated with the acquisition team, and the effect of time to first contact cannot be reliably assessed from this dataset.

Assumptions I made

- Leads are counted in the month they were created, not the month they booked. Calls run on to 10 July, so counting the other way would distort the later months.
- Where I averaged attempts per lead, I counted leads with no calls as zero rather than leaving them out.
- Bookings come from the prospects table rather than the booking flag, because the two disagreed before I removed the duplicates.
- For the conversion window check I used 27 June as the cutoff. That is 10 July minus 13 days, the longest gap between a lead and its booking that I found.

What the data cannot answer

- Everything stops at the booked visit. There is nothing on whether the visit happened, whether it turned into a customer, what the contract was worth, or how long people stay. Without that I cannot say which channel is really the best one.
- I cannot explain why bookings per connected lead fell, which is where most of the decline sits. There is no information on targeting, creative, bidding or competitors.
- There is no agent cost or capacity data, so I cannot work out whether the extra calls are worth making.
- Leads that got later calls got them because someone decided to keep calling, so those rates describe what happened rather than what an extra call would achieve.